

ECE421: Lab 1 Q&A
Bryan Yan (yanbryan) and Lyu Wang (wanglyu)

Part 1b

Answer the following question(s), write and save your answer in a separate PA1_qa .pdf file. Remember to submit this file together with your code.

1. Refer to the documentation, what is the functionality of the `tol` parameter in the Perceptron class? (2 marks)
2. If we set `max_iter=5000` and `tol=1e-3` (the rest as default), does this guarantee that the algorithm will pass over the training data 5000 times? If not, which parameters (and values) should we set to ensure that the algorithm will pass over the training data 5000 times? (2 marks)
3. How can we set the weights of the model to a certain value? (2 marks)
4. How close is the performance (through confusion matrix) of your NumPy implementation in comparison to the existing modules in the `scikit-learn` library? (2 marks)

1. Tol is the tolerance parameter, if it is not None, then training will stop once the improvement in error (loss function) is smaller than the tol value. I.e., updates to the weight vector will stop on the iteration where `loss > previous_loss - tol`
2. It does not guarantee the algorithm will pass over the training data 5000 times. The `max_iter` makes training stop at or before 5000 epochs, and tol also is a condition to end training earlier if the loss isn't improving by more than tol.
3. To manually set the model's weights, we can simply assign numpy arrays to the `model.coef_` and `model.intercept_` (the former sets the weights, the latter sets the bias terms). To force the Perceptron to always train for exactly 5000 iterations, we should disable early stopping by setting: `Perceptron(max_iter=5000, tol=None)` Setting `tol=None` ensures that the model never stops early based on convergence criteria.

Now, the Perceptron will always iterate through the training data for the full `max_iter=5000`.

4. It varies, but generally the results from our implementation are fairly close to the `scikit-learn` implementation

```
(.venv) C:\UofT\ece421\ece421-labs\PA1>python PerceptronImp.py
-----Test Result-----
Confusion Matrix is from Part 1a is: [[6. 0.]
 [5. 9.]]
Confusion Matrix from Part 1b is: [[ 6  0]
 [ 2 12]]

(.venv) C:\UofT\ece421\ece421-labs\PA1>python PerceptronImp.py
-----Test Result-----
Confusion Matrix is from Part 1a is: [[12. 0.]
 [ 0. 8.]]
Confusion Matrix from Part 1b is: [[12  0]
 [ 0  8]]

(.venv) C:\UofT\ece421\ece421-labs\PA1>python PerceptronImp.py
-----Test Result-----
Confusion Matrix is from Part 1a is: [[ 8. 0.]
 [ 1. 11.]]
Confusion Matrix from Part 1b is: [[ 8  0]
 [ 0 12]]
```

Part 2a

Answer the following question(s), write and save your answer in a separate PA1_qa.pdf file. Remember to submit this file together with your code.

- When we input a singular matrix, the function `linalg.inv` often returns an error message. In your `fit_LinRegr(X_train, y_train)` implementation, is your input to the function `linalg.inv` a singular matrix? Explain why. (2 marks)
- As you are using `linalg.inv` for matrix inversion, report the output message when running the function `subtestFn()`. We note that inputting a singular matrix to `linalg.inv` sometimes does not yield an error due to numerical issue. (1 marks)
- Replace the function `linalg.inv` with `linalg.pinv`, you should get the model's weight and the "NO ERROR" message after running the function `subtestFn()`. Explain the difference between `linalg.inv` and `linalg.pinv`, and report the model's weight. (2 marks)

1. The input to `linalg.inv` (which we later changed to `linalg.pinv`) is `X_train.T @ X_train`, which may be a singular matrix (specifically it will be singular if the columns of `X_train` are not linearly independent).
2. I get "ERROR", because `linalg.inv` can't compute the inverse of a singular matrix, as its inverse just doesn't exist. But sometimes the parameter could be invertible, so `linalg.inv` does not report an error.
3. I now get "NO ERROR", with the weight shown below. Unlike `inv`, `pinv` computes the pseudoinverse, which for nonsingular matrices matches the normal inverse, but for singular matrices, also exists.

`linalg.inv(A)`: Computes the exact inverse of a matrix only if it is nonsingular (full-rank). If `A` is singular, an error occurs.

`linalg.pinv(A)`: Computes the Moore-Penrose pseudo-inverse, which is defined for all matrices, including singular and non-square matrices.

For nonsingular matrices, `pinv` gives the same result as `inv`.

For singular matrices, `pinv` provides the closest possible solution, making it useful in least squares regression.

```
(.venv) C:\UofT\ece421\ece421-labs\PA1>python LinearRegressionImp.py
-----subtestFn-----
weights: [1.04360964e-14 2.00000000e-01 4.00000000e-01]
NO ERROR
-----testFn_Part2-----
Mean squared error from Part 2a is 3046.582974183576
Mean squared error from Part 2b is 3046.5829741835773
```